8

Distinguishing

The previous chapter discussed one way to shore up inference about a mechanism—identifying as many implications as possible and assessing them in as many ways as possible. But it is more than just the number of implications and assessments that matters. Suppose two models, each embodying different mechanisms, each generates the same implications. Then the elaboration cannot adjudicate among the mechanisms. Assessing common implications only helps assess whether at least one of the mechanisms is at work. Any stronger claim is open to immediate reinterpretation.

How can we better assess whether a particular mechanism is at work? We can identify and evaluate implications that are not common implications of the alternatives on offer. We call such an implication a *distinguishing implication.*

Many mechanisms are at work in most real-world targets. So it's important not to overinterpret a claim about distinguishing. Distinguishing is rarely about arguing for one mechanism and against others. Rather, it is about increasing how much we know about whether a particular mechanism is at work. In some circumstances, good news for one mechanism may indeed be bad news for another. But distinguishing is essentially about the question: "How convincing is the case that this mechanism plays a role in explaining what's going on in the target?" This is an importantly different question from: "Which mechanism explains what's going on in the target?"

To get a feel for how distinguishing works, let's return once again to Dell's (2015) work on the Mexican drug war. She shows an association between PAN election victories and increases in drug-related violence. Her interpretation points to increased crackdowns by the PAN. These crackdowns weakened incumbent drug traffickers, creating opportunities for other

organizations to try to wrest control of territory through violence. But other mechanisms might be at work as well. To take just one example, the mere fact of a change in political power might lead to violence. Recall that the PAN was mostly the challenger party. So PAN victories are associated with the instability that accompanies a change in political control. As an alternative mechanism, this could be the explanation for increased violence in places where the PAN won the mayoral election.

Dell identifies an implication of her interpretation that distinguishes it from this alternative. On her account, violence should increase more when there is a crackdown than when there is not. Only the PAN implemented a crackdown policy. Thus we should expect an increase in violence when a PAN challenger wins a close election over a non-PAN incumbent, but not when a non-PAN challenger wins a close election over a PAN incumbent. The alternative account implies an increase in both cases because both cases result in a change in political control. Dell compares the change in violence following turnovers in which PAN mayors displace incumbents of other parties to those in which PAN incumbents themselves are displaced. She finds that violence increases when the party in control switches to the PAN but not when it switches away from the PAN. This bolsters her claim that the crackdown mechanism is at work.

In this chapter, we discuss several papers that distinguish between alternative models. As was the case for elaborating, distinguishing can take many forms. After describing the papers, we discuss some strategies that are specific to distinguishing over and above elaborating.

8.1 Accountability and Term Limits

The family of accountability models introduced in Chapter 7 represents two mechanisms: electoral incentives and electoral selection. An incumbent who will stand for reelection shapes her actions with an eye toward convincing the voter she is a good type. And candidates who have won past elections differ systematically from those who have not had to face a process that selectively filters out lower types.

The electoral incentives mechanism is intimately related to classical questions about accountability from democratic theory. As such, distinguishing the electoral incentives mechanism from the electoral selection mechanism has become particularly salient in the literature.

To distinguish the electoral incentives mechanism, we need a way to compare politicians with different electoral incentives. A first thought is that term limits provide that comparison. Incumbents who are term limited cannot stand for reelection and so have weaker electoral incentives than incumbents who are not term limited and can stand for reelection.

Unfortunately, most comparisons of term-limited and non-term-limited incumbents fail to hold all else equal. Imagine comparing senators, who are never term limited, to governors in a state that has gubernatorial term limits. Too much differs between the two offices for a reasonable interlocutor to accept that all else has been held equal. Similarly, comparing incumbent legislators in states with and without legislative term limits would not provide a credible estimate of electoral incentive effects. Too much else differs across states.

The obvious solution is to compare term-limited and non-term-limited incumbents in the same political office, in the same political system. For example, a scholar could study a legislature with term limits, and compare members for whom the limit did and did not bind. This comparison does hold equal significantly more than the governor–senator or cross-state comparisons. But it does not hold electoral selection equal. To see the problem, imagine a simple setting in which no one can serve more than two terms. Electoral incentives imply that first-term incumbents should behave differently than second-term incumbents. But electoral selection implies that second-term incumbents are, on average, better than first-term incumbents. So differences in performance between term-limited and non-term-limited incumbents are a common implication of both mechanisms.

There are ways forward here. The two papers we discuss next try, in different ways, to combine term limits with a research design that also holds equal electoral selection.

8.1.1 Mayors in Brazil

Ferraz and Finan (2011) study the electoral incentives mechanism in the context of Brazilian mayors. For measurement, they exploit an anticorruption policy implemented in 2003. Starting that summer, the central government used lotteries to select municipalities for audits. These lotteries were staggered, with fifty or sixty municipalities audited in most months. All of these audits covered the period from 2001 to 2003. The program was unanticipated,

so the actions taken prior to the announcement of the audits were unlikely to be affected by the possibility of random audits.

Ferraz and Finan use the corruption uncovered by these audits as a measure of whether incumbent politicians take costly actions that benefit voters. They compare the corruption of first-term mayors (who were eligible for reelection) to that of second-term mayors (who were not eligible for reelection). All else equal, the incentives created by electoral accountability imply that first-term mayors should engage in less corruption.

The model that Ferraz and Finan have in mind to motivate their study is as follows. Suppose there are corruptible and incorruptible types. The corruptible types make up a share $\lambda \in (0, 1)$ of politicians, while the incorruptible types make up a $1 - \lambda$ share. In each period, corruptible types choose to engage in an amount $a \in [0, 1]$ of corruption, while incorruptible types do not engage in corruption. The voter observes a noisy signal of corruption.

Suppose the equilibrium level of corruption by corruptible types in their first term in office is $a^* \in (0, 1)$ and the equilibrium level of corruption by corruptible types in their second term is 1, reflecting the absence of electoral incentives in the second period. Moreover, electoral selection implies that the share of corruptible types among second-term incumbents is $\underline{\lambda} < \lambda$.

A comparison of corruption by second-term mayors to corruption by first-term mayors includes both electoral selection and electoral incentives. The expected level of corruption in the first period is λa^* and the expected amount of corruption in the second period is $\underline{\lambda}$. The electoral selection component is reflected in the fact that second-term mayors are less likely to be corruptible types ($\underline{\lambda} < \lambda$). The electoral incentives component is reflected in the fact that, conditional on being a corruptible type, second-term incumbents engage in more corruption than do first-term incumbents ($1 > a^*$).

Ferraz and Finan estimate that second-term incumbents divert between 2 and 3 percentage points more resources to corruption than do first-term incumbents.

As we've already noted, this difference does not isolate the effect of electoral incentives. If the electoral selection mechanism works as in the model, some of the difference reflects the fact that second-term incumbents are, on average, less corruptible than first-term incumbents. As such, this finding may be an underestimate of the true effect of electoral incentives.

But it is also possible that selection works in the opposite direction. For instance, corruptible incumbents might be more willing to engage in corrupt

election practices, and that could lead to positive selection for corruption. So the simple comparison does not distinguish the mechanisms as clearly as we would like.

To address this worry, Ferraz and Finan compare second-term incumbents to just those first-term incumbents who will go on to win reelection. All of these incumbents, first- and second-termers alike, won reelection at the conclusion of their first term. As such, the distribution of types is likely more similar between the first- and second-term mayors in this restricted sample than in the simple comparison of all first- and second-term mayors.

To see this in the model, note that the share of corruptible types among first-term incumbents who win reelection is the same as among second-term incumbents, $\underline{\lambda}$. But they differ in their effort due to differing electoral incentives. A comparison of the amount of corruption by second-term incumbents and first-term winners yields the quantity $\underline{\lambda}(1 - a^*)$. This quantity reflects the effect of electoral incentives $(1 - a^*)$ weighted by the share of corruptible types among those who win reelection $(\underline{\lambda})$.

The empirical analysis finds that first-term incumbents who go on to win reelection behave 3–4 percentage points less corruptly than do second-term incumbents. This estimate is consistent with our argument, earlier, that the simple comparison of second- to first-term mayors underestimated the incentive effect. Moreover, it supports the distinguishing implication, bolstering the claim that electoral incentives are at work.

8.1.2 State Legislatures

Fouirnaies and Hall (2018) study electoral incentives in the context of US state legislatures. They collect data on bill sponsorship, roll-call voting, and committee assignments, covering legislatures in fourteen states that have legislative term limits. They use principal components analysis to construct a one-dimensional measure of legislative productivity out of these data.

Their main analysis regresses productivity on an indicator variable for being term limited, with individual legislator and chamber-by-year fixed effects. Intuitively, this research design compares, for members of a particular legislature, the change in productivity from year t to year $t + 1$ for those who are term limited in year $t + 1$ to the change for those who are not term limited.

A virtue of this difference-in-differences type design is that it controls for fixed legislator attributes that may be the target of electoral selection. For example, suppose that each legislator's competence stays the same over

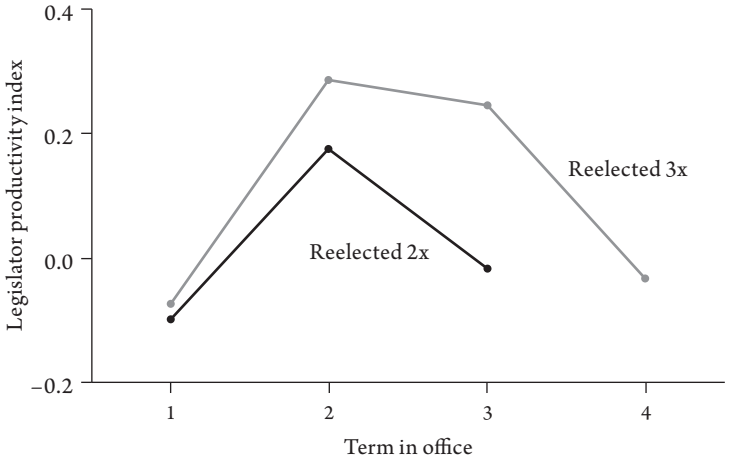

FIGURE 8.1. Term-limited legislators are less productive in their final term. Based on Fouirnaies and Hall (2018), Figure 1.

time. Electoral selection implies that the legislators who make it to their final allowed term have a better distribution of competences than do those who lose an election along the way. But, as long as an individual's competence isn't changing systematically over time, the fixed effects account for this difference, because the analysis is within-legislator. In other words, Fouirnaies and Hall are asking whether the same legislator behaves differently when she is able to run again or not.

Figure 8.1 illustrates the design. The gray curve represents the productivity over time for legislators who left office via term limits. The black curve represents the productivity over time for legislators who do not win reelection into the last term allowed by the term limit. Both graphs show a decline between the penultimate and final terms, but the decline is greater in the gray curve. Fouirnaies and Hall interpret this to mean that term limited, but not other, incumbents lose electoral incentives for effort.

Although the design controls for differences in incumbent characteristics that are fixed over time, we still might worry about Fouirnaies and Hall's comparison of term-limited incumbents to incumbents who have won fewer elections. Even without differences in legislator characteristics, this comparison may not hold all else equal.

One example, which Fouirnaies and Hall emphasize, is that there may be learning by doing in the legislature, and the magnitude of that learning may change over a career. It is plausible, for example, that legislators learn less

between their third and fourth terms than between their first and second. If this is the case, then Fouirnaies and Hall's comparison overstates the effect of electoral incentives.

To partially address such concerns, Fouirnaies and Hall replicate their analysis with cohort-year fixed effects in place of chamber-year fixed effects. This means that they compare changes in productivity for legislators who have served the same number of terms, some who face binding term limits and others who do not. Although this analysis now makes comparisons between states, if state legislatures that differ in the number of terms they allow are otherwise comparable, then learning on the job should be similar for term-limited and non-term-limited incumbents. The results of this analysis are very similar to the primary specification.

Another concern is that the results might be explained by reversion to the mean in legislative performance. To examine this possibility, Fouirnaies and Hall elaborate two additional implications of the electoral incentives mechanism that are not subject to concerns about reversion to the mean. First, incentives should be stronger the higher is the salary of a member of the legislature. Second, incentives for term-limited legislators are not completely absent if they can run again after a term out of office. And, indeed, Fouirnaies and Hall find that states with higher legislative salaries or lifetime bans on returning to the legislature have larger declines in productivity for term-limited legislators. These patterns are implications of the electoral incentives mechanism, but do not seem to be implications of learning or of reversion to the mean, bolstering the case that the incentives mechanism is at work.

8.2 Party Effects in Congress

In Chapter 6, we left the debate over the role of parties versus ideology in determining legislative behavior at something of an impasse. Rohde (1991) argued that over-time increases in party-unity voting suggest parties matter. But Krehbiel showed that such increases are a common implication of two different mechanisms, one based on party influence and the other on ideological sorting.

Krehbiel (1993) also pointed out a distinguishing implication of the party influence mechanism. If there is party influence ($\beta > 0$ in the model from Chapter 6), then two congresspeople with the same ideal point, but belonging to different parties, will vote differently on some bills. By contrast, if there is no party influence, these two members will vote the same way on every bill. If

there is overlap in the distribution of ideal points, so that there are congresspeople with the same ideal points who belong to different parties, it is possible to empirically evaluate this distinguishing implication.

The ideological voting mechanism also has an implication that distinguishes it from the party influence mechanism. Ideological voting implies that two congresspeople in the same party, but who have different ideologies, will vote differently on some bills, even though both face the same party pressure.

Ansolabehere, Snyder, and Stewart (2001) set out to assess these distinguishing implications. A first challenge in doing so is measurement validity. The theory makes a distinction between the preferences of the members and the positions they take in response to party pressure. So it is crucial that the measure of ideological preferences not be tainted by party influence. That means measures based on roll-call votes (e.g., NOMINATE) are inappropriate. Instead, Ansolabehere, Snyder, and Stewart use data from a survey called the National Political Awareness Test (NPAT). The NPAT asked all congressional candidates in 1994, 1996, and 1998 a battery of questions on a wide range of topical political issues. Fifty-seven percent of all members of Congress completed the NPAT at least once.

Are the NPAT scores a good measure of ideology, in the sense of being similar to the ideal points in our model? In general, a member's vote depends on many factors—party pressure, the member's own personal beliefs, the policies desired by constituents, pressure from donors, and so on. In the model, the ideal point stands in for all of these, except for party pressure. From this perspective, it is okay that the responses to the NPAT are unlikely to be pure reflections of the respondent's personal beliefs. Indeed, we don't want them to be. Instead, measurement validity rests on two key claims. First, NPAT scores are less likely to be affected by party pressure than is roll-call voting. Second, since the NPAT responses were public, the respondents would reasonably want to shape their answers to please their voters, donors, and so on. As such it might be reasonable to treat the NPAT answers as similar to the ideal points in the model. But if the NPAT responses reflect party influence or do not reflect constituent or donor demands, but only the representative's personal ideology, then the analysis will not succeed in holding all else equal in the comparison of Democrats and Republicans.

Ansolabehere, Snyder, and Stewart distill the NPAT answers into a single ideological score for each candidate using a simple scaling technique. Comparing roll-call voting records to the NPAT-based ideological scores, they can assess the extent to which Republicans and Democrats with similar measured ideologies vote differently. This allows an assessment of the

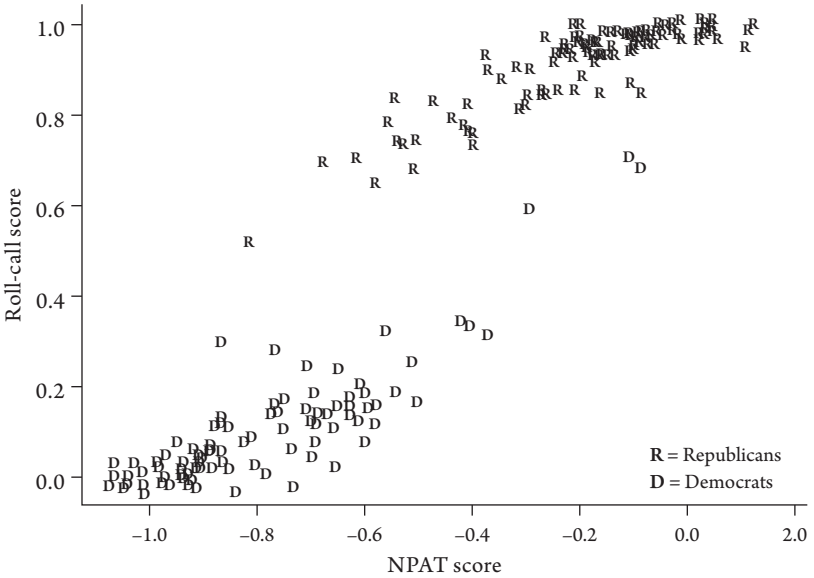

FIGURE 8.2. Replication of Ansolabehere, Snyder, and Stewart's (2001) Figure 2, showing the ideology and roll-call voting behavior of individual Democratic and Republican congressional representatives for the 105th Congress.

distinguishing implication of the party influence mechanism—representatives of different parties, but with the same ideology, will sometimes vote differently.

Figure 8.2 shows Ansolabehere, Snyder, and Stewart's analysis in a nutshell. The horizontal axis represents the scaled NPAT ideology scores, while the vertical axis represents scaled roll-call voting scores. In other words, the horizontal axis represents the candidates' stated ideological positions, while the vertical axis represents how they actually voted in Congress. The "D"s and "R"s on the graph represent individual Democrats and Republicans.

There are Democrats and Republicans with similar ideological positions in the middle range of the horizontal axis. The gap between the Democrats and Republicans in this region reflects the difference in voting records between members of different parties, but similar scaled NPAT responses. To the extent that we believe the NPAT score is a reasonable measure of ideology in our model, the gap shows that ideological sorting cannot account for the data on its own—Democrats and Republicans with the same ideology vote differently on average. This agrees with the distinguishing implication of the party influence mechanism.

At the same time, the data show that voting records of members of each party are positively correlated with ideology. This evidence is consistent with the distinguishing implication of the ideological mechanism, although this is not the authors' main focus.

While it takes an important step forward in the literature, the Ansolabehere, Snyder, and Stewart study is not without its own limitations. Candidates may respond to voter influence differently on the NPAT and in actual roll-call voting. If candidates attempt to appear, say, moderate on the NPAT, but then vote their true ideology once in Congress, then the results of Figure 8.2 could arise without any party influence. On the other hand, suppose that candidates who are more moderate than the party overall attempt to appear to be in line with the party by giving NPAT answers that reflect the standard party platform even when they disagree with it. Then Figure 8.2 would underestimate the extent of party influence because party influence would already be at play in the NPAT scores themselves.

A second measurement concern is that the roll-call voting measure, NOMINATE, may not capture the same issues that are covered in the NPAT survey. If the issues that received roll-call votes were more polarizing than the NPAT questions, then the differences evident in Figure 8.2 could simply reflect the differing nature of the issues considered rather than party influence. In principle, it should be possible to match NPAT questions to roll-call votes on the same issues, but to our knowledge this has not been done.

Measurement concerns aside, let's take stock of the progression from Rohde to Krehbiel to Ansolabehere, Snyder, and Stewart. Rohde claims that increased party-line voting indicates that party influence has increased in Congress over time. Krehbiel reinterprets, suggesting that evidence of increased party-line voting is consistent with a purely ideological mechanism and increased ideological sorting. But Krehbiel doesn't only reinterpret. He also suggests a distinguishing implication of the party influence mechanism. Evidence agreeing with that implication would constitute more compelling evidence of party effects. Ansolabehere, Snyder, and Stewart do not engage the question about change over time. Instead, they pick up Krehbiel's distinguishing challenge and attempt to provide evidence about party-line voting not confounded by ideology. By reducing the plausibility of the claim that the evidence can be entirely explained by ideological differences between Democrats and Republicans, Ansolabehere, Snyder, and Stewart make it harder for a reasonable interlocutor to claim that parties don't matter for roll-call voting in Congress.

But this is not the end of the story. Members of the same party do not all vote the same way on every roll call. We have identified several factors that could account for this heterogeneity—personal preferences, constituent or donor demands, and so on. Although this relationship is not Ansolabehere, Snyder, and Stewart's main focus, a future researcher with a compelling research design might be able to assess the importance of these more fine-grained mechanisms.

8.3 Social Pressure and Political Behavior

An important strand of literature on political behavior contrasts two different visions of what motivates people to act politically. To what extent are people motivated by their intrinsic individual interests or values and to what extent are they motivated by the desire to conform to extrinsic social norms? To start approaching this question, it is useful to be able to distinguish the two different theoretical mechanisms.

Empirically assessing the influence of these mechanisms on political behavior in a way that holds all else equal is challenging. Suppose individuals who vote express a higher sense of civic duty, for instance in a survey. It would not be obvious whether the relationship is causal or whether there are other factors—perhaps related to an individual's upbringing—that lead to both a sense of duty and a propensity to vote. Similarly, if people in the same social network were observed to vote for similar candidates, the similarity could be due to social pressure imposed by the network or it could be the result of homophily, the tendency for birds of a feather to flock together. Given these concerns about unobservable confounders, a just-controlling research design is unlikely to uncover credible evidence about the all-else-equal implications of the intrinsic or extrinsic mechanisms in question.

As such, experiments have come to play an important role in this literature. Here we discuss two contributions, both of which use experiments to assess distinguishing implications of the social pressure mechanism. One focuses on voter turnout and the other on racial solidarity in politics.

8.3.1 Social Pressure and Voting

The question of why people vote has received considerable attention from both theorists and empiricists (Feddersen, 2004; Blais, 2006). A leading explanation is that people vote, at least in part, out a sense of civic duty (see

Riker and Ordeshook, 1968, for an early formulation). Gerber, Green, and Larimer (2008) set out to distinguish two possible mechanisms that might underlie such a sense of duty. The first is intrinsic: the reward that the voter obtains from the act of voting, or the guilt she would feel from sitting at home on election day. The second is extrinsic: the reward that the voter obtains when others observe her voting, or the shame she would feel from others learning that she failed to perform her civic duty. A priori, it is plausible that both mechanisms are at work.

Gerber, Green, and Larimer conduct a large-scale experiment in which some participants received a treatment that is meant to activate the intrinsic sense of duty to vote, and other participants received one of several other treatments that primed both intrinsic and extrinsic considerations to varying degrees. The setting is in the weeks prior to the August 2006 statewide primary election in Michigan. The authors sent mailings to 80,000 households, divided into four treatment groups. (Another 100,000 households received no mailing and constituted an untreated group.) All of the treated households received postcards that reminded them of the date of the upcoming election. The first treatment group's postcards primed the intrinsic duty to vote with a reminder to "DO YOUR CIVIC DUTY AND VOTE" [capitalization in original]. The three remaining treatment groups received postcards that included that intrinsic duty prime along with an additional message meant to prime feelings of social pressure to vote, with each consecutive treatment meant to produce increasing degrees of social pressure. One group received a postcard noting that voting is a matter of public record and "YOU ARE BEING STUDIED"; another group received a postcard displaying the voting records of those within the household and reading "WHO VOTES IS PUBLIC INFORMATION"; and the final group received a mailing displaying both the household's turnout history and neighbors' turnout history, with the question "WHAT IF YOUR NEIGHBORS KNEW WHETHER YOU VOTED?" These last two treatments indicated that a follow-up mailing after the election would report to either the household or the neighborhood the subject's turnout in the upcoming election. Importantly, each of the three social pressure treatment groups also received the message to "DO YOUR CIVIC DUTY—VOTE!" So we can think of the social pressure mailings as priming the combined sense of intrinsic and extrinsic motivations.

How does Gerber, Green, and Larimer's experiment distinguish extrinsic from intrinsic motivations for voting? Comparing the first treatment group to the untreated group provides evidence about the combination of the

TABLE 8.1. Voting Rates Under Various Treatments from Gerber, Green, and Larimer's (2008) Experiment on Social Pressure and Voting

	Experimental Group				
	No Mailer	Civic Duty	Being Studied	Voting Public	Neighbors
Percent voting	29.7	31.5	32.2	34.5	37.8
Observations	191,243	38,218	38,204	38,218	38,201

reminder of the election and intrinsic motivations, since it only primes the sense of civic duty. The remaining three treatments prime various extrinsic considerations in addition to reminding voters about the election and priming intrinsic motivations. Thus, an increase in turnout for any of these three treatment groups—relative to the stand-alone civic duty treatment—is an implication of the extrinsic mechanism but not the intrinsic mechanism. In addition, because each of the three remaining treatments are meant to prime an increasingly strong sense of extrinsic motivation, an increase in turnout across the three remaining treatment groups is also an implication of the extrinsic mechanism but not the intrinsic mechanism. Thus, we can see Gerber, Green, and Larimer as assessing two distinguishing implications of the extrinsic mechanism.

Table 8.1 shows the turnout rates Gerber, Green, and Larimer report for each treatment group. In additional analyses not reproduced here, the authors show that all of the pairwise differences between groups are statistically significant. Each treated group turned out at a higher rate than the untreated. Individuals in the civic duty treatment were roughly 2 percentage points more likely to vote than those in the untreated group, and the treatment effects increased with the strength of the social pressure message for the three remaining groups. Of particular note, individuals in the fourth treatment group, who saw both their own and their neighbors' turnout history, had the highest turnout rate and were roughly 8 percentage points more likely to vote than those in the untreated group. This is a very large treatment effect in the context of the get-out-the-vote literature (see Gerber, Green, and Larimer, 2008, Table 4).

The results of the experiment provide support for two implications that distinguish the extrinsic mechanism from the intrinsic mechanism: turnout is higher for all of the social pressure treatment groups relative to the civic duty treatment group, and turnout rates increase with the intensity of the

social pressure treatments. Taken together, these results provide support for the claim that the extrinsic mechanism is at work in explaining voter turnout.

8.3.2 Social Pressure and Racial Solidarity

Questions about the role of social norms are particularly important in the context of the political behavior of Black Americans. As White, Laird, and Allen (2014) argue, Black Americans' political behavior is highly uniform even though the considerable economic and religious heterogeneity in the Black community might suggest substantial divergence of individual policy interests. Is this because Black Americans, while primarily motivated by private interest, use group interest as an informative proxy for private interest (Dawson, 1995)? Or is it because social pressure and social norms, in addition to private interests, play a major role in shaping political behavior?

White, Laird, and Allen seek to distinguish the private interest mechanism and the social pressure mechanism in an experiment on college students who attend a historically Black college or university (HBCU). The particular behavior they study is political donations during the 2012 presidential election between Barack Obama and Mitt Romney.

They divide students into an untreated group and two treated groups. Each member of the untreated group is simply asked to choose a division of $100 into an amount to be donated to Obama and an amount to be donated to Romney. While no actual political donations get made, the study participants are not made aware of this until the end of the experiment.

Members of the first treated group are given this same task. But they are also told that the computer will randomly give them an incentive—identifying one of the candidates and offering them $1 in personal compensation for every $10 that they donate to that candidate. In actuality, the computer always identifies Romney as the incentivized candidate. The idea is to induce private incentives to support the candidate opposed by the social norms of the student's community at the HBCU. Unlike the political donations, subjects actually receive these personal incentive payments.

Members of the second treated group are given the same task and the same incentives as the members of the first treated group. But, in addition, they are told that their names and donations will be published in the campus newspaper. The idea here is to induce students to consider the possibility that they will face social sanctions for behaving against social norms.

TABLE 8.2. Donations to Obama Under Various Treatments from White, Laird, and Allen's (2014) Experiment

	Experimental Group		
	No Incentive	Incentive	Incentive and Newspaper
Average Obama donation	$90.22	$68.10	$85.98
Observations	48	50	50

How does this experiment distinguish the private interests from the social norms mechanism? An implication of the private interests mechanism is that, all else equal, a person with a private interest in donating to one candidate should be more inclined to donate to that candidate. The social norms mechanism has no such implication. So, the first treated group donating more to Romney and less to Obama than the untreated group is a distinguishing implication of the private interests mechanism. Similarly, an implication of the social norms mechanism is that, all else equal, a person whose actions will be made public should behave more in conformity with social norms than a person whose actions will remain private. The private interest mechanism has no such implication. So, the second treated group donating less to Romney and more to Obama than the first treated group is a distinguishing implication of the social norms mechanism.

White, Laird, and Allen's findings are reported in Table 8.2 (which is based on their Figure 3). The table shows the average donation to Obama for each of the three groups. (Students had to donate the full amount, so donations to Romney are the balance.) Students in the control group donated an average of about $90 to Obama. Students in the group that were incentivized to donate to Romney decreased their donations to Obama to about $68. This suggests the private interest mechanism is at work. Making donation behavior public increased donations to Obama back up to $86, almost the unincentivized level. This suggests the social norms mechanism is also at work.

8.4 Gleanings

Distinguishing requires elaborating. So all the approaches described in Chapter 7 are relevant here. This chapter illustrates some additional strategies that are particularly useful for distinguishing. All of them boil down to controlling away the effects of other mechanisms to isolate the mechanism

in question. But there are some interesting nuances that are worth drawing attention to.

One approach is to construct comparison groups in a way that plausibly holds constant the effects of other mechanisms. To get a distinguishing implication of the party influence mechanism, Ansolabehere, Snyder, and Stewart compare voting records of representatives with the same ideology but different party affiliations. The studies by Ferraz and Finan and by Fouirnaies and Hall compare politicians with the same degree of electoral selection, but who face different electoral incentives.

An experimental setting opens up additional possibilities, since the researcher, at least in part, constructs the environment being studied. Gerber, Green, and Larimer's field experiment and White, Laird, and Allen's lab experiment both exploit this opportunity by designing experiments with multiple treatments. In each case, one treatment is designed to activate one mechanism, and the other treatment adds to the initial design a feature that activates the second mechanism. Comparing the two distinguishes the second mechanism from the first.