Optional Essays — Persuasion Theory and the Interpretation of Survey Experiments

Jacob Edenhofer*

06/05/2026

Purpose

These essays are aimed at students whose research design proposals are organised around an empirical contribution — in particular, a survey experiment, an information-provision experiment, or the analysis of an existing one — rather than around the development of a formal model. We covered the relevant material only briefly in Week 1, so the questions below do a fair amount of hand-holding, including a recap of the persuasion-theory background you will draw on.

There are two reasons why an empirically-minded student should work through one of these essays. First, the persuasion-theory recap below introduces a small piece of formal-theory notation — senders and receivers, types θ, messages m, posteriors $\Pr(\theta \mid m)$, the distinction between the prior and the posterior — that recurs in many other (game-theoretic) models you are likely to encounter. Becoming comfortable with it on a familiar applied case (interpreting a survey experiment) makes the same notation easier to read elsewhere. Second, working through a published experiment with this framework forces you to identify the assumptions that are doing the analytical work — which mechanism makes the message credible, which stage of the response is being measured, what would have to be true for the average effect to be commensurable with the theoretical prediction — and that skill is exactly what the main assignment asks you to apply to your own argument.

*jacob.edenhofer@nuffield.ox.ac.uk

What to Submit

Choose one of the two tasks below and email a single document of approximately 900–1,200 words. You may attempt both if you wish; in that case the combined submission should not exceed 1,800 words. You may use the section headings provided, or merge them as you see fit. The worked half-example at the end of this document shows what one paragraph of completed analysis should look like.

The two tasks differ in the kind of experimental design they ask you to interpret:

Task A An *information-provision* experiment, in which a treatment group receives a factual claim and the outcome is some belief or attitude.

Task B A *priming* or *cue-taking* experiment, in which a treatment group is exposed to a stimulus or source label and the outcome is an attitude or stated preference, with no new factual content provided.

For each task I provide a suggested paper, but you are welcome to use another paper of similar standing if it fits your proposal better.

Background You Will Need

The three subsections below provide the persuasion-theory background you will draw on. The framework follows Little (2023), *Bayesian Explanations for Persuasion*, supplemented at the end by the three-stage decomposition standard in the empirical literature on elite cues and information provision.

The puzzle of persuasion

The starting point is an impossibility result. Suppose a sender (politician, pundit, researcher, interest group) wants a receiver (voter, respondent, citizen) to take a particular action — to vote a certain way, to update a belief, to support a policy. Suppose the receiver is Bayesian, i.e. she updates her beliefs by Bayes' rule on whatever message she receives. Then, by the law of iterated expectations, the receiver's average posterior must equal her prior. The sender cannot, on average, move beliefs.[1]

[1] Formally, let $\theta \in \{0, 1\}$ denote the underlying state, with $\theta = 1$ corresponding to good news and $\theta = 0$ to bad news, and let $p \equiv \Pr(\theta = 1)$ denote the receiver's prior. The sender observes θ and chooses a message m according to a strategy $m^*(\theta)$; the receiver observes m and forms the posterior $\Pr(\theta = 1 \mid m)$ by Bayes'

The intuition is straightforward. Imagine a message that, in equilibrium, made the receiver more confident that the news was good. Knowing this, a sender who simply wants the receiver to think the news is good would send that message regardless of the underlying truth. But if the sender always sends the same message, the receiver extracts no information from it, and it ceases to be persuasive. Persuasive equilibria therefore unravel when a Bayesian receiver, a sender with monotone preferences over the receiver's action, and no further structure.

This is why the question "how is persuasion possible at all?" is non- trivial. The five mechanisms below are five distinct ways of breaking the impossibility result.

Five mechanisms

1. Costly signalling. The message itself costs something to send, and — crucially — the cost is higher for the sender type with bad news than for the sender type with good news. This is known as the *single-crossing* condition, and it makes the costly message a credible signal of type, in the tradition of Spence (1973) and Milgrom and Roberts (1986).

The classical illustration is the labour market. An MBA is costly to obtain, but more costly for a candidate of lower ability, who has to study harder for each course to pass. The degree teaches little that is useful for the job, yet employers can still infer ability from who chose to obtain it. Only candidates whose underlying ability is high enough that the

rule. The claim is that, averaging across the messages the receiver expects to see,

$$\mathbb{E}_m\big[\Pr(\theta = 1 \mid m)\big] \;=\; \sum_m \Pr(m)\,\Pr(\theta = 1 \mid m) \;=\; \sum_m \Pr(\theta = 1, m) \;=\; \Pr(\theta = 1) \;=\; p.$$

The first equality is the definition of expectation; the second uses Bayes' rule to rewrite each posterior as a joint probability ($\Pr(m)\Pr(\theta = 1 \mid m) = \Pr(\theta = 1, m)$); the third marginalises over m. This is sometimes called the *martingale property* of Bayesian belief updating. The intuition behind it is the following. If the receiver knew, ex ante, that her belief would on average rise upon hearing the sender's message, she would already have revised her prior upward in anticipation of receiving it; revising afterwards would be a free correction she had failed to make in advance. Beliefs formed by Bayes' rule are by construction the receiver's best current guess given everything she knows, including her model of the sender's strategy, and the best current guess cannot be expected to move on average without already being wrong. Colloquially: you cannot, on average, be surprised. Combined with a sender who simply wants the receiver to take a higher action, the martingale property delivers an impossibility. Suppose a candidate equilibrium had the receiver's posterior strictly above her prior upon observing $m = 1$, and hence, by the equation above, strictly below her prior upon observing $m = 0$. A sender holding bad news ($\theta = 0$) then has a strict incentive to deviate and send $m = 1$ regardless, because doing so induces the higher receiver action she prefers. Once both sender types send $m = 1$, the receiver extracts no information from the message, and the posterior on observing $m = 1$ collapses back to p. The candidate informative equilibrium unravels. This is the content of Theorems 1 and 2 in Little (2023). The five mechanisms in the body text are five distinct ways of breaking the unravelling argument: by making messages differentially costly to send across types, by aligning sender and receiver preferences, by ruling out outright lying, by giving the sender a stake in future interactions, or by allowing the sender to pre-commit to a messaging strategy before learning the state.

resulting wage premium exceeds the personal effort cost find the degree worth pursuing. The general structure is that *sending the message* is itself an action, and that action carries a cost differential across types large enough that, in equilibrium, only the good-news type bears it. The receiver can then take the message at face value.

Examples in politics include taking principled stands that alienate part of one's coalition (the cost is paid in lost donations or (primary) votes; only candidates whose information truly justifies the stance bear it); expensive campaign advertising (only campaigns expecting ad spending to translate into votes invest at scale); time-intensive volunteering and door-to-door canvassing; costly protest in repressive regimes (where the cost is paid in police violence or arrest); and the reputational cost of backing down after a public commitment. The key intuition is that people may engage in seemingly wasteful or even harmful behaviour because the wastefulness is what makes the signal credible. Only senders who really do hold the good information find it worthwhile to incur the cost.

Application to survey experiments. The mechanism is most useful as a *substantive* object that an experiment can manipulate. A vignette or prompt can describe an action that the relevant political actor took — a candidate's costly stand against her party's leadership; a campaign that spent heavily on door-to-door canvassing in unfavourable territory; an interest group that publicly burned bridges with one of its allies — and the design can vary the cost while holding the underlying type fixed. The treatment then isolates the inference respondents draw from the cost differential.

2. Common interests. Even when messages are not costly, communication can be partially informative if the sender's preferences are partially aligned with the receiver's, in the cheap-talk tradition of Crawford and Sobel (1982). The smaller the bias, the more information transmits in equilibrium; perfect alignment supports a fully honest equilibrium.

Example. Asking a friend to recommend a restaurant. Your friend wants you to enjoy the meal, you want to enjoy the meal, the friend has no reason to lie, and the recommendation is informative. Asking a friend's medical advice: again, mostly aligned — unless the friend works for a drug company that profits from a particular treatment, in which case the bias is large and the advice less reliable. The Crawford–Sobel result formalises this dependence on alignment. The state space gets partitioned into communicable categories, and the partition becomes finer (more information transmits) as the sender's bias relative to the receiver shrinks; below a threshold of alignment all messages are babble or "cheap talk".

A subtler implication, which Little emphasises, is that a sender who is biased on average can still be credible on dimensions where the bias is smaller — "of the things I want you to do, X matters more to me than Y" can be informative even when both X and Y serve the

sender's interest. A parent who tells a child "if you only have time for one subject tonight, do maths over English" is making a comparative claim that does not strictly serve the parent's interest in either direction taken alone, so revealing the ranking is plausible. The same logic applies to a lobbyist who tells a legislator "of the regulatory changes my industry wants, the licensing reform matters more to us than the tax credit". Both serve the lobbyist, but the ranking is informative.

Application to survey experiments. Common interests is most often the natural justification for why a factual claim in information-provision experiments from a researcher moves respondents at all. The researcher and the respondent share a truth-tracking goal in the experimental context, the bias is small, and the equilibrium supports honest reporting. Whether respondents *perceive* this alignment is itself testable by design. A vignette or prompt can vary the apparent alignment between source and respondent — a health claim attributed to "your local doctor" versus "a pharmaceutical company spokesperson"; a policy claim attributed to a co-partisan or out-partisan think tank — and measure whether respondents update more on the more aligned source. Heterogeneity in updating across these conditions is direct evidence on the common-interests mechanism. Partisans who treat an academic source as out-group will violate the alignment assumption even for the experimenter herself; subjects who suspect they are being persuaded by the researcher may partially discount the claim.

3. Verifiable information. The sender cannot lie – she can either reveal the truth or stay silent. This is sometimes called the hard-information or persuasion-game framework (Grossman, 1981; Milgrom, 1981). With *monotone receiver preferences* — meaning that the receiver's optimal action moves in the same direction as her belief about the state, so a more favourable posterior leads to a higher action and a less favourable posterior to a lower one — the unravelling argument applies. The type with the very best news reveals; the type just below has an incentive to reveal so as not to be pooled with everyone worse; the next type has the same incentive; and the chain runs all the way down. The result is full revelation.

A standard illustration. Imagine a job applicant who can choose to disclose his GPA or not. The applicant with a 4.0 strictly prefers to disclose. Knowing this, an employer reading "no GPA disclosed" should infer that the applicant's GPA is below 4.0, which sharpens the incentive of the 3.9 applicant to disclose, and so on. In equilibrium, every applicant above the very lowest possible GPA discloses because the alternative is being pooled with the lowest type. The same logic applies to politicians releasing tax returns, scientists preregistering studies, and companies disclosing financial statements. Silence becomes a signal of "I have something to hide," which forces revelation by every type for whom revealing is

even marginally better than being pooled with the worst.

Two non-obvious implications follow. First, even types with mediocre news will reveal it because staying silent is taken as worse news. So, we should not expect senders to share only favourable information. Second, the mechanism collapses if receivers fail to discount silence appropriately. Brundage et al. (2022) call this *selection neglect* and use it to explain why partisan media can durably shift average beliefs even though, in equilibrium, they should not.

Application to information-provision experiments. Verifiability is the natural credibility mechanism for information-provision designs. The corrective claim is sourced from official statistics or peer-reviewed evidence, in principle checkable, and respondents respond to the treatment as if the claim cannot simply be made up. The mechanism predicts a particular pattern of heterogeneity. Respondents with access to verifying evidence outside the survey — the highly informed, or those embedded in counter-messaging environments — should update less, because the experimental signal is one of several verifying signals available to them. Respondents who treat the source as plausibly fabricating (e.g., partisans who suspect academic sources of ideological bias) will also update less than the unravelling logic predicts. For priming experiments verifiability is largely irrelevant, because no factual claim is being made.

4. Reputation. The sender cares not only about the current decision but about being seen as competent or honest in future interactions. Reputation cuts both ways. A reputation for *competence* can support honesty when the sender wants to be seen as well-informed about the state, but it can equally induce *pandering* — saying what the audience already believes, in order to look smart, or claiming greater certainty than the underlying signal warrants, because confident-sounding senders look more competent. A reputation for *honesty* supports informativeness more straightforwardly but typically requires repeated interactions, so that today's lie has a future cost.

A useful illustration is the bond-rating agencies. Moody's and S&P have a long-run reputational stake in being accurate, because clients pay for ratings only if other market participants treat the ratings as informative. This pushes the agencies toward truthful disclosure. But the agencies are also paid by the issuers they rate, which pushes them toward favourable ratings in the short run. The infamous mid-2000s ratings of mortgage-backed securities are a case in which the short-run incentive to please clients overwhelmed the long-run reputational concern — and the long-run reputational concern was eventually validated when the agencies' credibility collapsed in the financial crisis. The same tension applies to cable-news pundits, financial advisers, and political commentators.

Application to survey experiments. Reputation is mostly a property of the *branded source* embedded in the treatment ("According to The New York Times…"), not of the researcher. A study that varies the source label across conditions is implicitly running a reputation experiment, and the analytical question is then which source has more reputation capital with which subgroups. A treatment that randomises whether a fact is attributed to a co-partisan or out-partisan source, holding the fact constant, is a clean reputation manipulation; the heterogeneity in the treatment effect tells us about subgroup-specific reputation rather than about the underlying factual claim.

5. Commitment (Bayesian persuasion). The sender pre-commits to a messaging strategy *before* observing the state, and the receiver knows the strategy. Persuasion then operates by choosing the strategy that maximises the sender's expected utility ex ante, which often means lying as much as possible subject to keeping the receiver willing to follow the signal. The framework was popularised by Kamenica and Gentzkow (2011) and has since become dominant in the formal-theory literature.

The mechanism is best illustrated by their canonical example. A prosecutor (sender) wants a judge (receiver) to convict a defendant. The state is either "innocent" or "guilty," with some prior probability of guilt below one-half — so absent any signal, the judge would acquit. The prosecutor designs an investigative procedure (a forensic test) and is bound to its results before the defendant's true type is observed. Suppose the procedure always returns "guilty" when the defendant is guilty, and returns "guilty" with probability μ when the defendant is innocent. The prosecutor chooses μ in advance. The judge convicts on a "guilty" verdict if her posterior belief in guilt exceeds one-half. The prosecutor's optimal strategy is the largest μ consistent with that constraint — that is, she lies as often as possible while keeping the receiver's incentive constraint just satisfied. Persuasion comes from the sender's ability to bind herself to a strategy that, ex post, she would want to deviate from; the receiver, knowing the binding holds, treats the message as informative.

Plausible real-world examples are typically institutional. Editorial guidelines at a media outlet — a durable, ex ante choice about how often to criticise a government, made before any individual story arises — can be approximated by a commitment strategy because revising the guideline for each new story is costly. Commissioning a study with a pre-registered design has the same structure – the design choices are locked in before the data arrive.

The cleanest case of a real-world commitment device is, ironically, a pre-registered survey experiment itself. The experimenter writes the survey instrument, the treatment text, and the analysis plan in advance. Once the study is in the field, the experimenter cannot adjust the message based on what she learns about the respondent. This makes the commitment

framework relevant for thinking about the experimental *apparatus* — it is part of why pre-registration can help credibility that ex post communication cannot — but not, by itself, a persuasive model of how parties or pundits move voters in the wild.

The three-stage framework

My summary of the empirical literature on elite cues and information provision (see Druckman, 2022; Coppock, 2023) is that a treatment effect can be decomposed into three stages, with failures of persuasion being possible at each stage.

Reception The message reaches the receiver and is registered. Failure here means the respondent never noticed the prime, did not read the information carefully, or could not recall it minutes later.

Persuasion Conditional on reception, the message moves the receiver's beliefs (information-provision) or the salience of an existing consideration (priming). Failure here means the message was registered but did not change the relevant cognitive object: the prior was already at the truth, the consideration was already at maximum salience, or motivated reasoning blocked the update.

The Posterior and beyond
Conditional on persuasion, downstream attitudes, behaviour, and durability follow. Failure here means beliefs or salience changed but the attitude did not move (because the attitude rests on many beliefs of which the treated one is not pivotal), or the effect decayed within days, or it did not generalise to a related outcome.

A central empirical observation: many disagreements about whether "voters update on information" are really disagreements about *which stage* moved or failed. The framework's discipline is to keep the three apart.

1 Task A — Interpreting an Information-Provision Experiment

Suggested paper. Hopkins, Sides, and Citrin (2019), "The Muted Consequences of Correct Information about Immigration" (*Journal of Politics*). This paper is well-suited to the framework because it provides correct information about the size of the immigrant population, finds that beliefs do update, and finds that downstream attitudes move much less than

the magnitude of the belief update would suggest. You are welcome to use another paper (e.g. Alesina, Miano, and Stantcheva, 2023) if it fits your research interests better.

Setting up the case

In one paragraph, explain the factual claim that constitutes the treatment, the treatment-control contrast, and the outcomes measured. Then identify which of Little's five mechanisms the design *implicitly* invokes for the message to be informative. Defend this rather than asserting it: would all subgroups perceive the alignment? would all subgroups treat the source as verifiable?

Tracing the three stages in this paper

For your chosen paper, answer:

(a) Was reception measured by a manipulation or attention check, and what share of respondents passed?

(b) Was belief updating measured separately from downstream attitude change, and were both pre- and post-treatment beliefs collected?

(c) Was durability tested in a follow-up wave? Was a behavioural outcome included?

The point is to be precise about which stage the paper provides direct evidence on.

The commensurability problem at the persuasion stage

Information-provision experiments report an average treatment effect on beliefs. But Bayes' rule does not predict that every respondent updates by the same amount. A respondent whose prior already matches the truth receives a non-informative signal — the likelihood ratio is one — and rationally does not move. A respondent with a very precise prior moves only a little even when wrong. The ATE pools these respondents with those whose priors are far from the truth, and a null or modest average effect is not, by itself, evidence against rational updating: the experimental quantity and the theoretical quantity are not commensurable without conditioning on the prior.

A concrete illustration. Suppose the actual share of foreign-born residents is 14% and the experiment provides this figure to the treatment group. A respondent who already believed 14% should not update at all. A respondent who believed 30% should revise downward. A respondent who believed 5% should revise upward. The ATE on beliefs is a weighted

average of these heterogeneous predicted effects, with weights determined by the population distribution of priors. So a small ATE can be perfectly consistent with the persuasion mechanism working exactly as it should for the respondents on whom it is predicted to work.

Answer the following:

(a) Does the chosen paper measure priors and condition on them?

(b) If not, what does the average effect tell us, and what does it not tell us?

(c) Is there a subgroup of theoretically null-predicted respondents (those whose priors were already correct) who are nevertheless interesting to study — for example, because if they *do* move, that is potential evidence on motivated reasoning or experimenter demand rather than on Bayesian updating?

Examining alternative explanations

Pick the strongest pattern in the paper — "beliefs move, attitudes don't"; "no movement at all"; "heterogeneous effects by party" — and walk through the alternative explanations that the framework supplies:

(a) reception failure rather than persuasion failure;

(b) pre-existing counter-cues drowning out the treatment;

(c) many-belief attitudes (one fact need not be pivotal for the attitude);

(d) source-credibility heterogeneity rather than motivated reasoning;

(e) divergence vs. desensitisation in the sense of Little, Schnakenberg, and Turner (2022).

For each, say what evidence in the paper would let you rule it in or out.

Policy relevance: partial vs. general equilibrium

A well-run information-provision experiment estimates a *partial-equilibrium* effect: what happens to one respondent who is forced to receive the signal, holding the rest of the information environment fixed. Policy advice — "fact-checkers should publish corrections," "a party should run an information campaign on issue X" — is a claim about a *general-equilibrium* outcome, in which behaviour by other senders, gatekeepers, and receivers also

changes in response. Two considerations, in particular, should make you cautious about the move from one to the other.

Self-selection of (news) sources. In the experiment, exposure to the treatment is forced. In the wild, the respondents most likely to be *moved* by a corrective fact are also the ones least likely to *encounter* it, i.e. they consume different outlets, follow different feeds, and avoid contradicting messages. The forced-exposure design measures a treatment effect on a population that often does not exist outside the survey environment. Reception failure in the field is selectively concentrated on exactly the respondents whose persuasion-stage effect was largest under forced exposure.

Counter-framing and incentive distortion. Once a fact or correction is fielded at scale, opposing senders counter-frame, and the producers of the original misinformation respond strategically. In misinformation experiments specifically, scaling fact-checking can *raise* the equilibrium incentive to lie more aggressively, because lies that escape fact-checking become disproportionately valuable. The partial-equilibrium estimate then over-states the policy effect, sometimes substantially.

Answer: in the chosen paper, does the partial-equilibrium estimate translate into the policy claim the authors actually make? Where does the chain break?

2 Task B — Interpreting a Priming or Cue-Taking Experiment

Suggested paper. Barber and Pope (2019), "Does Party Trump Ideology? Disentangling Party and Ideology in America" (*American Political Science Review*). This paper is well-suited because it varies a co-partisan elite cue (Trump's stated position) across policy items and shows that respondents move in the direction of the cue even on items inconsistent with their stated ideology. It is a clean test of cue-taking under the common-interests mechanism. Alternatives include Broockman, Kaufman, and Lenz (2024) on heuristic projection or Barber and Pope (2024) on issue importance.

Setting up the case

Identify the prime or cue, and the outcome. Is the treatment an *endorsement* (cue-taking) or a *vignette or frame* (priming proper)? The distinction matters because different sub-questions follow:

- Cue-taking shifts beliefs about the desirability of an option by attaching a source label

to it.

- Priming shifts the salience of an existing consideration without communicating new factual information.

If the paper claims to manipulate beliefs but the design only manipulates salience, that is itself worth flagging.

Locating the design within Little's mechanisms

For a cue-taking design, the most natural mechanism are either common interests (the source is taken to share the receiver's preferences) or reputation (the source has a track record of being right or honest). For a priming design, none of Little's five mechanisms applies in a straightforward way because nothing factually new is communicated. Priming experiments are interventions on which considerations a respondent draws on when answering.

Heuristic projection

For a cue-taking design, ask whether subjects are assumed to know the source's true alignment. Broockman, Kaufman, and Lenz (2024) document *heuristic projection*: respondents often assume an interest group or politician shares their own views, even when the source is in fact ideologically misaligned with them. Could the main effect in the chosen paper be projection-driven rather than cue-driven? What design feature would distinguish the two?

Tracing the three stages

(a) Reception. Is there a recognition or memory check? Subtle primes often fail at this stage silently — the participant never noticed.

(b) Persuasion (here: belief shift if cue-taking, salience shift if priming). Is it measured directly, or inferred from downstream attitudes?

(c) Posterior: magnitude, durability, scope. Priming effects are typically small and short-lived (Coppock, 2023).

Common mis-interpretations

For at least two of the four misreadings below, say whether the chosen paper guards against the misreading:

(a) treating lab priming as a field-level effect;

(b) treating a one-shot effect as a stable real-world effect;

(c) inferring the absence of the underlying consideration from a null priming effect (a possible ceiling effect on baseline salience);

(d) treating heterogeneous effects as evidence of *type* heterogeneity, when they may instead reflect heterogeneous baseline salience of the primed consideration.

Policy relevance: partial vs. general equilibrium

A well-run cue-taking experiment estimates a *partial-equilibrium* effect: what happens when one respondent is exposed to a cue, holding the rest of the information environment fixed. Policy advice — "a party should use cue X," "a candidate should be endorsed by source S" — is a claim about a *general-equilibrium* outcome, in which other parties, sources, and audiences respond. Two considerations matter.

Self-selection of (news) sources. Forced exposure to a partisan endorsement in a survey is not how endorsements reach voters in the wild. Voters who would be most moved by an endorsement from source S are typically the voters least exposed to S, because they sort into other outlets. The treatment effect inside the survey is therefore concentrated on a population that, in the field, is least likely to encounter the cue.

Counter-framing. In equilibrium, the moment a party or interest group fields a cue, opposing actors deploy a counter-cue. The experimental estimate is the effect of the focal cue holding the rest of the environment fixed; the general-equilibrium effect under contested communication is typically attenuated, sometimes sharply (Druckman, 2022, on credible counter-frames).

Answer: does the paper's policy claim — "parties should/should not use cue X" — survive these two corrections? If not, what is the smaller, partial-equilibrium claim that it does support?

A Worked Half-Example: Hopkins, Sides, and Citrin (2019) Through the Persuasion Framework

The example below illustrates the kind of analysis Task A is asking for.

Setting up the case. The explanandum is American attitudes towards immigration. The treatment is a corrective factual claim, with respondents in the treatment group being apprised of the actual share of foreign-born residents in the United States, while control subjects receive no such information. The outcomes measured are (i) the respondent's stated belief about the immigrant share, and (ii) downstream immigration-policy preferences. The implicit credibility mechanism is a combination of common interests (the survey researcher and the respondent share a truth-tracking goal in the survey context) and verifiability (the figure is sourced from official statistics and in principle checkable). Whether *partisan* respondents perceive the alignment is, however, an open question. A Republican respondent may treat the source as out-group, in which case the common-interests assumption fails and the verifiability of the claim does the entire credibility work.

Tracing the three stages. Reception is measured. Respondents are asked, post-treatment, what figure they were shown, and the relevant manipulation check is reported. Persuasion at the belief stage is measured directly. Pre- and post-treatment immigrant-share beliefs are collected, and the treatment moves them substantially. The persuasion stage to attitudes is the one that most fragile since stated immigration-policy preferences move much less than the size of the belief update would suggest. There is no follow-up wave, eaning durability cannot be assessed from the paper itself. The same holds for downstream behavioural outcomes.

The commensurability problem at the persuasion stage. The main finding — beliefs move, attitudes barely move — is often quoted as a challenge to rationalist accounts of voter behaviour. The persuasion framework qualifies this in two ways.

First, the average belief update is a weighted average across heterogeneous priors. Respondents whose priors were already close to the truth should not move, and respondents with very precise or strongly held priors should move only modestly. The ATE on beliefs is therefore not directly commensurable with a theoretical prediction without conditioning on the prior.

Second, the muted attitude response is consistent with a standard many-belief model of attitude formation: the immigrant share is one belief among many that feed into a stance on immigration, and updating only one of them is not predicted to move the attitude by a commensurate amount.

The main result that "information does not change attitudes" is therefore better stated as "information changes one input belief but not the attitude." Whether the underlying preference is unmoved *by anything* is a different and harder question.

Examining alternative explanations. (Sketch only — left for you to complete on your chosen paper.) Two competing explanations of the muted attitude response would be (i) source-credibility heterogeneity (the corrective is partially discounted by partisan subgroups, in which case attitude movement is concentrated in subgroups that already trust the source) and (ii) divergence in the sense of Little, Schnakenberg, and Turner (2022) (partisan respondents draw opposite attitude conclusions from the same belief update, so the average masks two non-zero subgroup effects). Distinguishing these requires variance heterogeneity tests at the attitude stage; the paper provides some evidence on the first but limited evidence on the second.

Policy relevance. The partial-equilibrium reading is straightforward. Forced exposure to a corrective fact moves beliefs but not attitudes.

So, does this imply that there is no point in providing this type of information? Here general-equilibrium considerations become relevant. Two considerations, in particular, matter. First, in the field, the respondents most moved by the corrective are the same respondents least likely to encounter it, because they sort away from sources that publish such corrections; reception failure in the wild is selectively concentrated where partial-equilibrium movement was largest. Second, an information campaign at scale would invite counter-framing from opposed actors, which the experimental control condition cannot represent.

While the paper makes a useful contribution, it should not be interpreted as evidence for or against the provision of such information. A single factual correction is not enough to move stated attitudes by much. Whether a sustained, contested information environment would do so is a question the experiment is not designed to answer – but this is the question that activists or policymakers must answer.